




Applied Azure AI Solutions Driving Frontier Transformation Across Enterprises and Industries

Nguyễn L. (Alfred) – Microsoft ASEAN Solution Engineer
alfredluu@microsoft.com

Frontier transformation

The new wave of growth & innovation beyond efficiency





Employee Self-Service

The "front door" for workplace services. Answers policy questions, retrieves HR data, and completes tasks seamlessly across HR apps and tools like HCM, benefits, pay, and career growth.



Hiring & Onboarding

Multiple agents designed to improve candidate and hiring manager experience and reduce time to hire. Capabilities include job description and interview question generation, start date negotiation (autonomous planning with candidate and hiring manager), manager and new hire collaboration (autonomous 1:1 coordination).



HR Professional Support

Curates HR professional content to simplify cross-functional engagement, upskill HR teams, and increase efficiency and reduce manual work in data gathering:

- Q&A on Talent, Org Development & Strategy content
- Provide quick access to HR experts
- Retrieve HR analytics data in natural language



Talent & Workforce Optimization

Allows users to reason over their org and skills data and run through workforce planning scenarios with the agent by providing workforce constraints, strategic plans, and requirements to model potential workforce changes.



Policy & Content Management

Augments and automates the policy lifecycle management process to ensure policy and procedure content is up-to-date, coherent, and relevant across the organization.

Early key learnings

- Build trust. Start with broad use cases like policy Q&A.
- Transform processes. Begin at the task level for control, then scale with multiple agents.
- Accelerate development. Utilize a mix of subject matter experts and developers.
- Define clear workflows. Agents for business processes require a combination of clear process definition and design, workflow automation, and agentic AI.

Operations

The Next Frontier of Banking

Customer Engagement

Corporate banking

- Enable instant global B2B payments
- Deliver real-time analytics and trade finance solutions

Corporate banking

Relationship managers

Relationship managers

- Shift human capacity from low-value tasks
- Focus relationship managers on deeper customer engagement

Risk and compliance

- Perform real-time liquidity stress testing
- Embed real-time AML and KYC compliance into core processes

Risk and compliance

Proactive and continuous approvals

- Deliver dynamic loan and mortgage servicing
- Provide proactive, personalized customer assistance

Lending and mortgages

Dynamic and autonomous

- Transform bank accounts into intelligent liquidity engines
- Allocate funds dynamically in real time

Deposits and payments

Cyber and fraud prevention

A new threat landscape

- Provide real-time, proactive cyber and fraud protection
- Secure all communications with quantum-safe encryption

Autonomous engagement channel

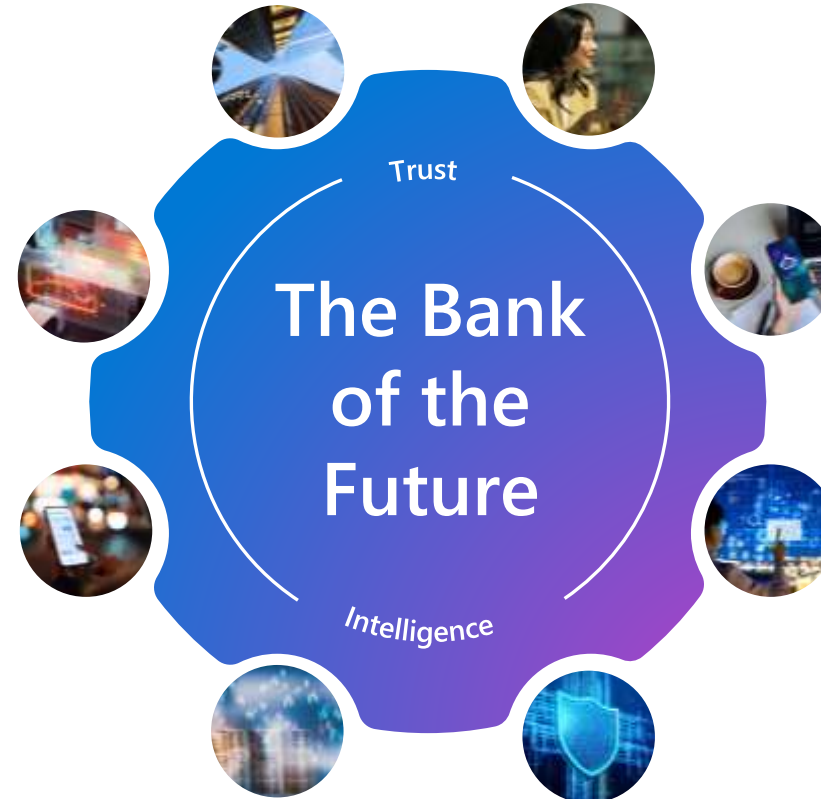
- Automate up to 95% of contact center interactions
- Deliver customer servicing through on-device AI assistants

Customer AI agent

A new channel

- Establish Know Your Agent (KYA) standards
- Enable autonomous agents to transact across customer segments

Agent to agent channel





Frontier transformation in Insurance



Accelerate underwriting

AI-driven risk modeling for instant, precise policy decisions



Evolve claims processing

Touchless claims processing with AI automating routine steps for fast payouts.



Transform customer service

Always-on AI assistants delivering personalized, omnichannel support.



Improve fraud detection

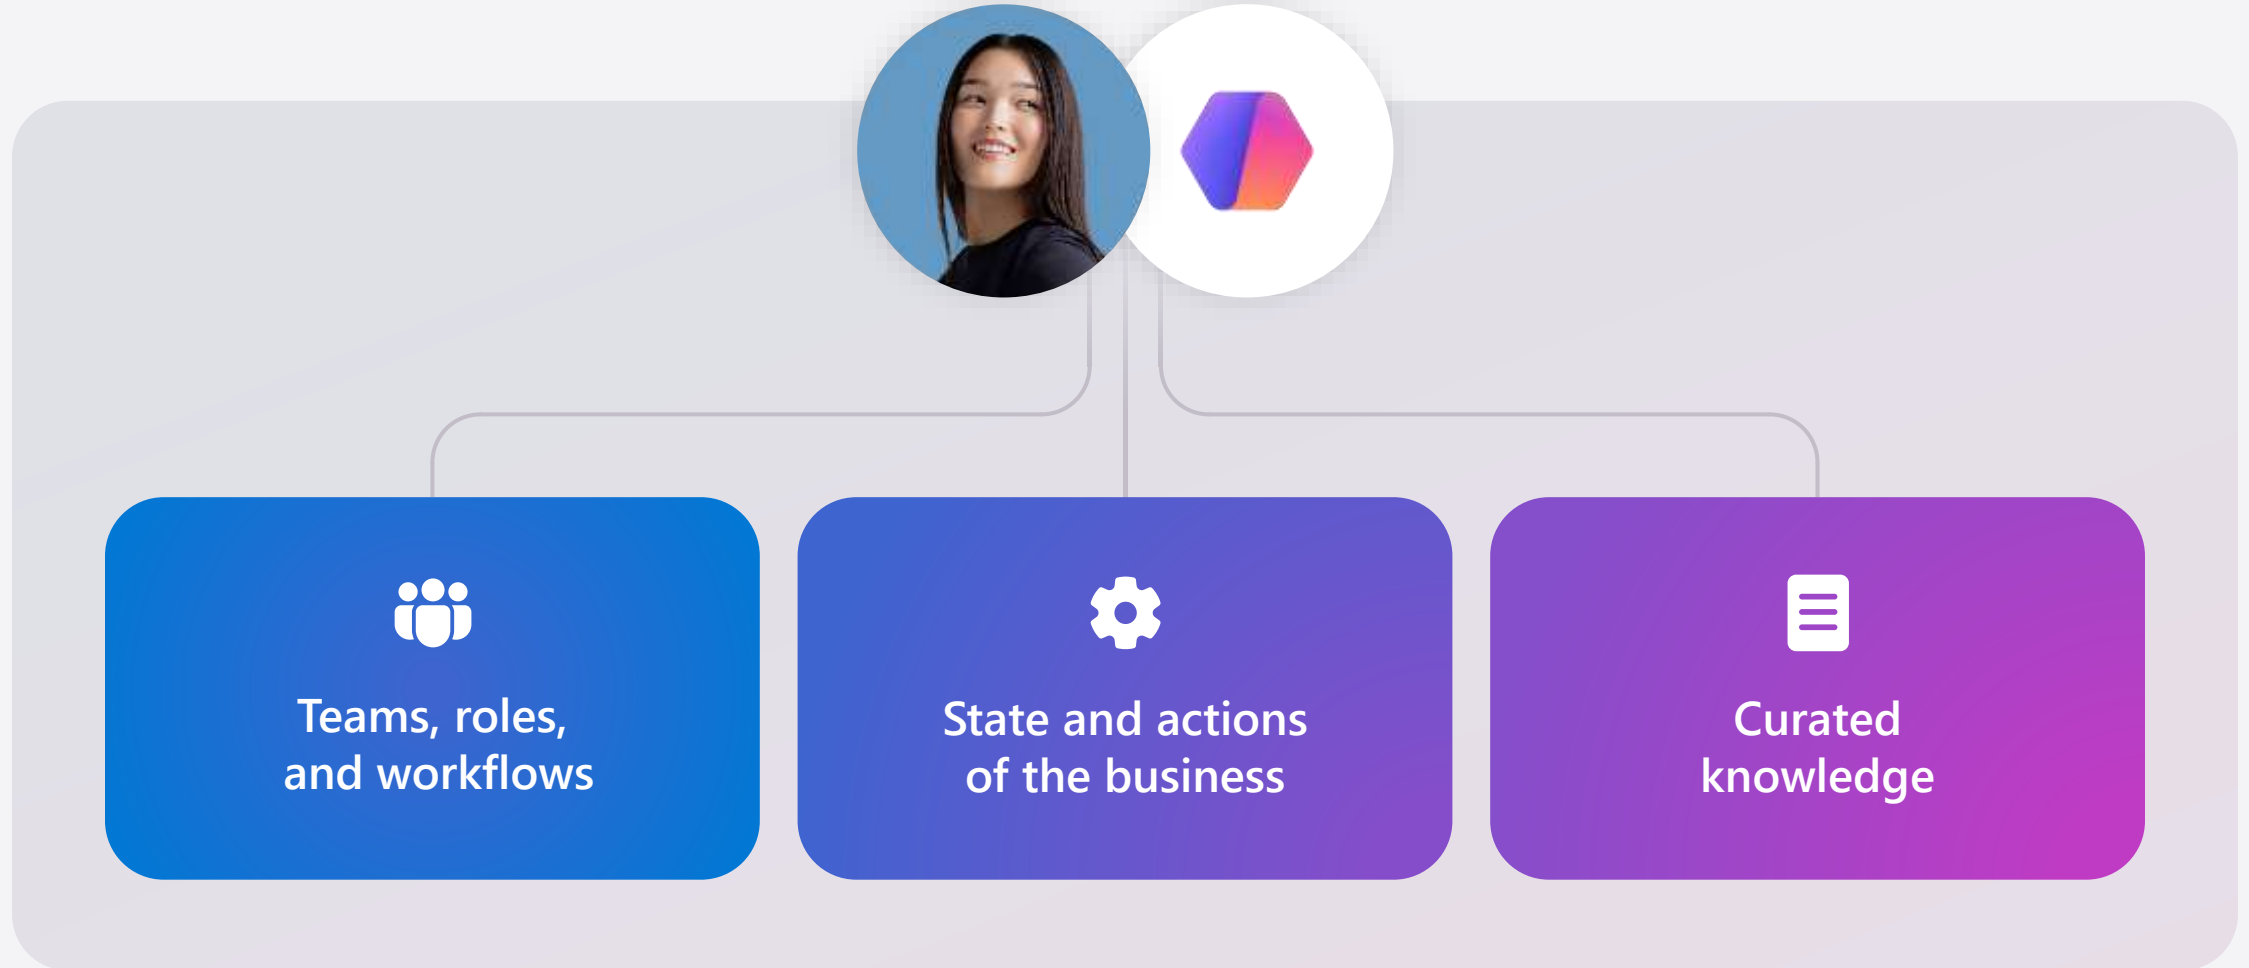
Real-time AI surveillance spotting fraud patterns and reducing false alarms.



Expand distribution

Data-driven product personalization and embedded insurance via AI insights.

What every **agent** needs to know



Your unique intelligence is what powers agents

Unified view
of your business

How your
employees work

Your knowledge
sources

Your business
processes

Your data and apps

Your IQ

Microsoft IQ

Unified intelligence for enterprise AI



Work IQ

"How your employees work"

Context on people, collaboration, and workflows



Fabric IQ

"How your business operates"

Context on business entities, systems of record, and actions

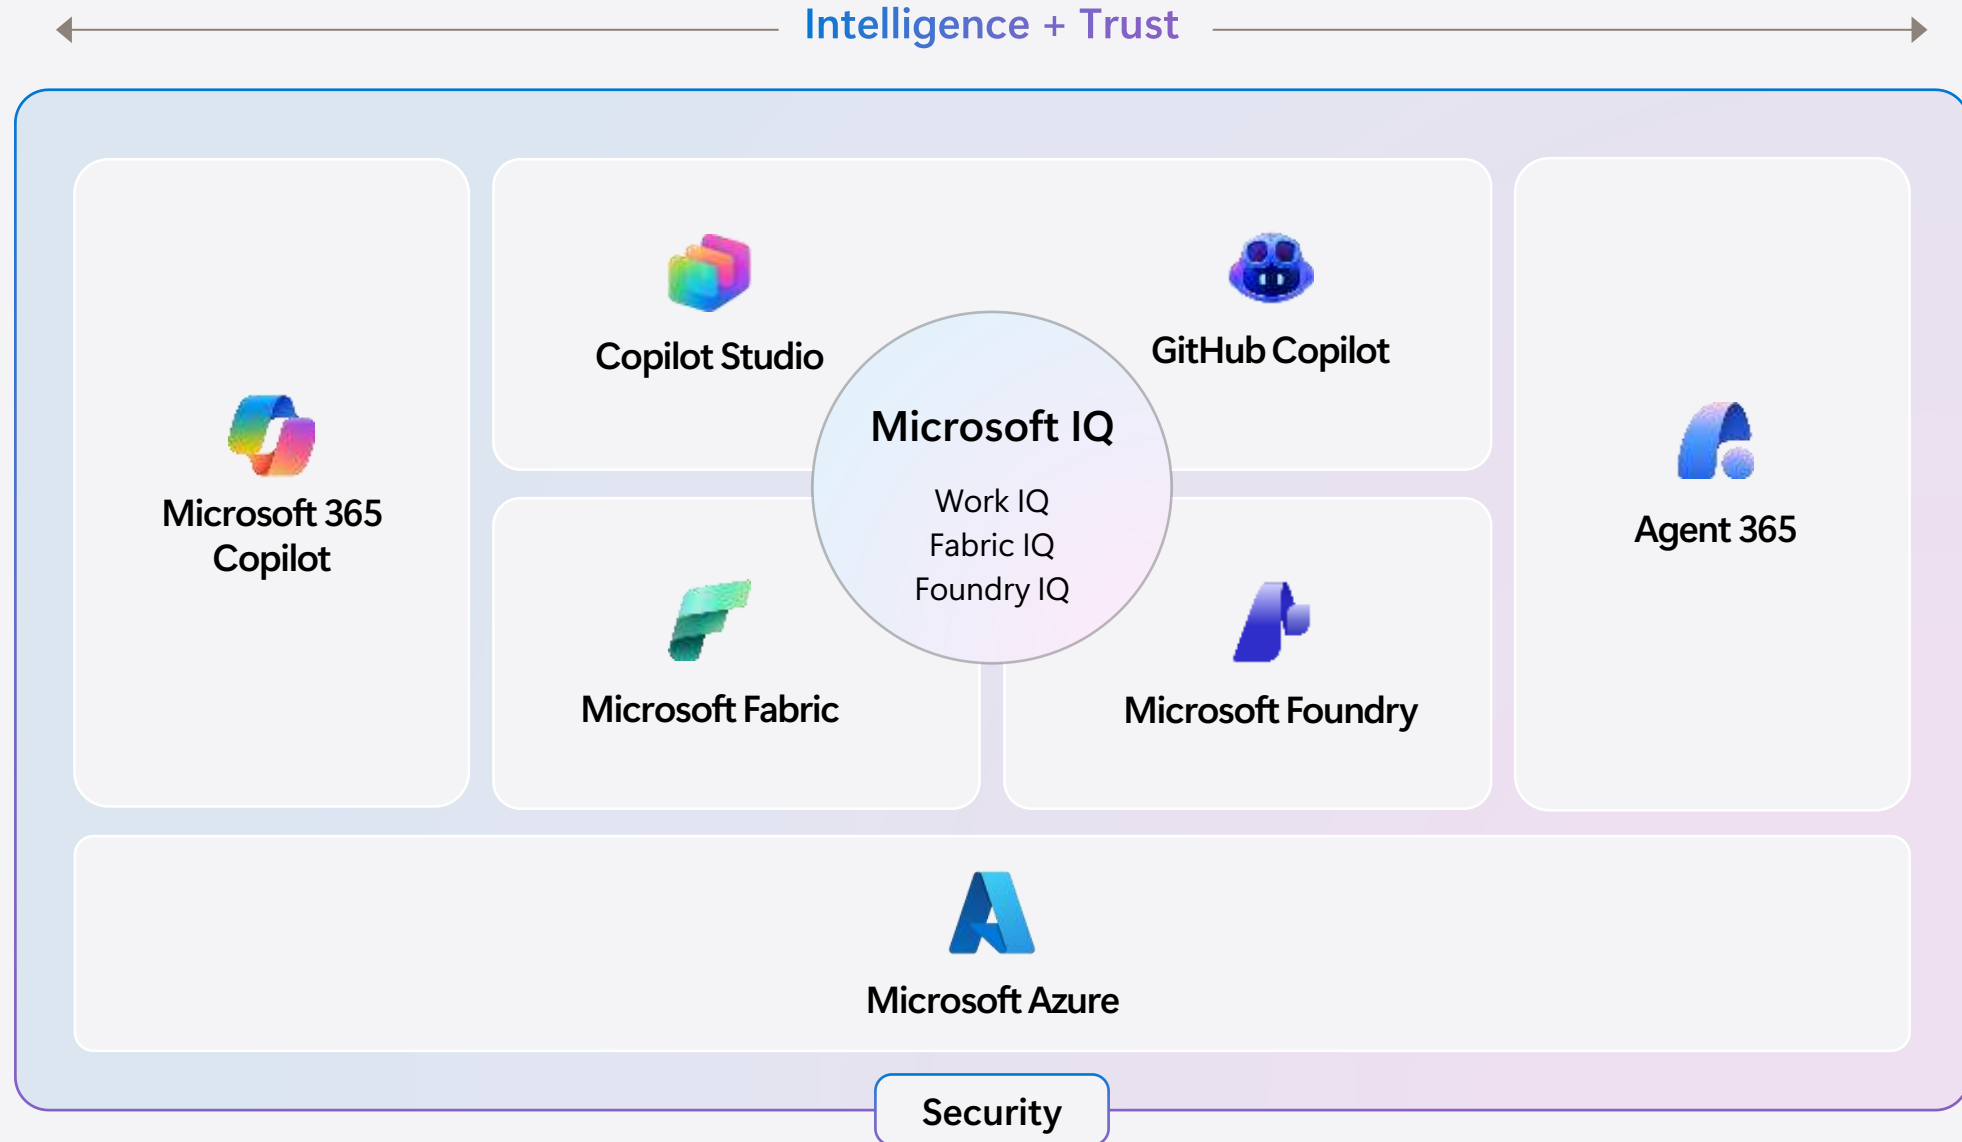


Foundry IQ

"How your agents unlock knowledge"

Context on policies, authoritative documents, and knowledge bases

Microsoft IQ Platform



Foundry Tools



80,000+ organizations using Microsoft Foundry today





Microsoft Foundry

The AI app and agent factory



Models



Agent Service



IQ



Tools



Machine Learning



Control Plane

Cloud

Edge

Security, compliance, and governance

Agents need contextual, trusted tools

Tools amplify intelligence



Context

Policies, data, and priorities

Lack of business context



Workflows

Repeatable and measurable

Manual workflows
slow productivity



Adaptable

Business needs and costs

High customization cost

Prebuilt, customizable tools, operate on your data, integrate with your business systems—
accelerating time to value.

The real-world hurdles of driving results with AI

Tool-calling is critical for agentic apps

Agents lack
business context
and don't drive
outcomes

Manual workflows
slow productivity

High
customization
cost slows
adoption

Organizations need prebuilt, customizable, action-oriented
tools that integrate with their business systems


Tools catalog



Introducing

Tool Catalog in Foundry


1,000+ tools
from public
tool catalog
from Microsoft
and partners




Build and
share tools via
organizational
private tool
catalog

aka.ms/FoundryToolCatalog

Empower agents to act on your business context



Agents can take a simple request and chain multiple tool calls together to complete a task





Azure Speech in Foundry Tools

Revolutionizing voice experiences with agentic AI

aka.ms/azure-speech

Azure Speech product overview

Speech to Text MCP

- Real-time transcription
- Fast transcription
- Batch transcription
- Custom speech

Text to Speech MCP

- Pre-built neural & HD voices
- Custom voices
- Text to Speech avatar
 - Video or Photo Public Preview

Speech Translation

- Real-time speech translation
- Live interpreter GA
- Video translation

LLM Speech Public Preview

- LLM powered Transcription and Translation

More tasks coming soon

Voice Live GA

- Real-time API for building voice agents
- Choice of models
 - Built-in
 - Bring your own

Voice agents are in demand, now more than ever



Customer service

Customer support and shopping assistance: voice chatbots for self-serve support, voice agent for call center and navigating product catalogs



Automotive

In-car voice assistant for command and control, general Q&A, etc.



Learning/education

Voice-enabled learning companion and training assistant



Human resources

Human resources support, such as employee support, career growth and development, job-related learning, interview assistant etc.



Public services

Voice agent for public services, such as citizen experiences, public administration assistance, etc.



Personal assistant

Voice-enabled agents that provide personal assistance to consumers, such as chatting, setting up reminders, managing calendars, etc.

Voice is the new interface between humans and computers

Challenges of building voice agents



Accuracy

High accuracy audio input and transcripts with customization for domain-specific vocabulary



Naturalness

Natural conversations with multi-lingual support and context awareness



Brand customization

Customized audio output for brand, image, and personality consistency



Flexibility

Choice of GenAI model—including both LLMs and SLMs



Robustness

Robustness of interruptions and end of turn detection in a variety of scenarios and noise environments



Immersive experiences

Visual identity for improved engagement and natural interaction

Engineering cost, time to market, and latency

Generally available

Voice Live API

Voice-enable any agent with advanced, real-time voice

Foundation model
choice

High quality global
locale coverage

More engaging with
Avatar

Easy customization

aka.ms/voice-live

Voice Live API

GA | Oct 2025

A single unified API with streamlined access to advanced speech capabilities and foundation model of your choice



Broad locale coverage and natural voices

140+ locales for speech to text (new) and 600+ prebuilt voices across 150+ locales for text to speech + Neural HD V2 voices optimized for conversations



Choice of Generative AI models

Option to choose from various gen AI models natively supported for your specific needs, including GPT-Realtime, GPT-5, GPT-4.1 series, Phi series and more; and to bring your own Foundry model (in preview)



Customization

Custom speech, custom voice, custom avatar (with 4K support), and lightweight speech input and output customization for higher accuracy



Agent integration, function calling

Seamlessly add voice input and output capabilities to agents built using Foundry Agent Service, improved function calls with tools



Conversational enhancements

Noise suppression, echo cancellation, robust interruption detection, etc. for natural conversations



Avatars

Easy-to-configure and avatar add-on to provide a visual identity for your voice agent



Call center integration

Easy to build call center voice agents with telephony channel integrated (through Azure Communication Service, and more to come)



Rich development resources

Support for C# and Python on top of the WebSocket protocol to simplify development, with additional programming languages to come

aka.ms/VoiceLiveGA

LLM Speech API

LLM enhanced speech transcription & translation through single API

Meeting notes



Voicemails



Captions & subtitles



Accurate,
context-aware &
multi-lingual

Add-ons including
diarization, wordtiming,
multi-channels

Prompt tuning

Ultra-fast inference

aka.ms/llm-speech

Public Preview

MAI-Transcribe-1

on Microsoft Foundry

~50% lower GPU cost than
leading alternatives

Enterprise-grade accuracy
across 25 languages

Powering Copilot Voice
Mode and Audio
Expressions

aka.ms/mai-transcribe-1-foundrycard



AZURE SPEECH



Public Preview

MAI-Voice-1

on Microsoft Foundry

Generates 60 seconds of expressive audio in under 1 second on a single GPU

Designed for highly expressive, human-like speech

Prompt the model with a 10-120s voice clip of your choice for instant voice cloning

aka.ms/mai-voice-1-foundrycard

GA | Ignite 2025

Live Interpreter API

Human-parity real-time speech-to-speech translation

Online meetings



Phone calls



Translator apps



E-Sports streaming



Multi-lingual

Automated language detection
of 76 languages

Significant latency improvement

reduces speech-to-speech translation
from 7s to 5s

Personal voice

preserves unique tone and style
of the original speakers

aka.ms/live-Interpreter

Demo In-car Voice Assistant





Azure Content Understanding in Foundry Tools

*Harnessing AI to extract structured insights from multimodal data —
accelerating decisions, reducing costs, and unlocking new revenue streams.*

Generally Availability (GA)



Azure Content Understanding

in Foundry Tools

Prebuilt &
custom analyzers


Multimodal content
processing

Validation with
confidence &
grounding

aka.ms/Foundry-Content-Understanding

Content Understanding Framework

Introducing a new bring-your-own-model deployment



← Try prebuilt analyzer

Document

Contracts

[Learn more about prebuilt analyzer](#)

Drag & drop files
here or

[Browse for files](#)

Sample



contract.pdf

Run analysis

Auto run analysis



WEB HOSTING AGREEMENT

This web Hosting Agreement is entered as of this 15 day of October, 2022 ("Effective Date") by and between **Contoso Corporation**, a Washington corporation, having its principal place of business at 1 Microsoft Way, Redmond, Washington, 98058, "Contoso" and **AdventureWorks Cycles**, a Washington corporation, having its principal place of business at 98 NW 76st Street, Suite 54, Bellevue, Washington, 98007, "AdventureWorks".

This agreement shall void and nullify any and all previous agreements to this date between Contoso and AdventureWorks.

There shall be no additional fees of any kind paid to Contoso, other than those stated within this agreement for software usage and/or bandwidth usage. AdventureWorks agrees to pay Contoso \$0.01 (one cent) per access up to 400,000 accesses thereafter payment shall be \$0.005 (one-half cent) per access. AdventureWorks shall send this amount to Contoso by no later than Wednesday for accesses used from the previous week (Monday thru Sunday).

Contoso must provide a person(s) to correct any technical problems (Server being down or inaccessible) 24 hours per day, 7 days per week. This person(s) must be available by beeper or telephone. AdventureWorks shall provide this same 24 hour service at the broadcast location.

This Agreement shall be governed by and construed in accordance with the internal laws of the State of Washington, applicable to agreements made and to be performed entirely within such state.

All parties have read and fully agree to all terms and conditions as set forth in this Web Hosting Agreement.

Contoso Corporation

By: Angel Brown
Title: CTO

Angel Brown

Adventure Works Cycle

By: Aaron Smith
Title: CEO

Aaron Smith

Fields

Result

Hide missing fields



● Title p.1

70.30%

WEB HOSTING AGREEMENT

● Parties (2) p.1

1 Name

67.20%

Contoso Corporation,

Address

69.80%

1 Microsoft Way, Redmond, Washington, 98058

ReferenceName

65.80%

("Contoso")

Clause

72.30%

Contoso Corporation, a Washington corporation, having its principal place of business at 1 Microsoft Way, Redmond, Washington, 98058 ("Contoso")

2 Name

74.30%

AdventureWorks Cycles,

Address

79.90%

98 NW 76st Street, Suite 54, Bellevue, Washington, 98007


Check token usage in JSON (Result tab) and learn about pricing [here](#).

Confidence scores

Grounded results

Tables and complex fields

Document segmentation and routing



Content Understanding Extraction Capabilities



Documents

- OCR – printed or handwritten text from 100+ languages
- Content elements (selection mark, check box, radio buttons, barcodes, formulas, image elements)
- Layout elements (pages, paragraphs, lines, sections)
- Tabular data (including complex multi-page layouts)
- Prebuilt Analyzers (including RAG-ready extractions)
- Multi-document input (Pro mode)
- Multi-file reasoning (Pro mode)
- Reference data (Pro mode)
- Classification & splitting (Pro mode)



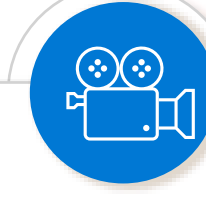
Image

- Generate titles, descriptions, and summaries
- Classify elements
- Prebuilt Analyzer



Audio

- Transcription
 - Diarization
 - Multi-lingual Transcription
 - Language Detection
 - Multi-Speaker role detection
- Generate conversation summaries and transcripts
- Determine sentiment
- Prebuilt Analyzers



Video

- Transcription
 - Diarization
 - Multi-lingual Transcription
 - Language Detection
 - Key frame extraction
 - Shot detection
- Segmentation (whole-video, automatic, custom)
- Face Detection and Recognition (gated)
- Generate summaries
- Determine sentiment
- Prebuilt Analyzers
 - RAG-ready extraction (prebuilt video indexing)

[What's new in Content Understanding? -| Microsoft Learn](#)

[Azure Content Understanding video overview - Microsoft Learn](#)



Azure AI Document Intelligence

Harnessing AI to accurately extract structure and schema—fueling automation, discovery, and faster, smarter decisions at scale.

Azure AI Document Intelligence



Document Analysis

Read

Extract printed and handwritten text

Layout

Extract text, tables and document structure



Prebuilt Models

Financial Services, Legal
(6 models)

US Tax (5 models)

US Mortgage (5 models)

Personal Identification
3 models



Custom

Custom Classification

Custom Extraction

Neural (recommended)

Template



Add on Capabilities

OCR (High Resolution, Formula, Font, Barcode)

Searchable PDFs
(only on **Read** model)

Key Value Pairs

Regulatory
Compliance



HITRUST



Choose the right service for document processing

Consideration Criteria	Azure Document Intelligence	Azure Content Understanding	Build Your Own with Foundry Models
Best Suited for	Structured/semi-structured documents like forms, invoices, IDs	Complex, high-variation, multimodal documents (text, image, audio, video)	Custom, niche workflows requiring full control over models and prompts
Input Types	Documents (PDF, images)	Documents, images, audio, video	Any, but requires preprocessing
Field Extraction	Prebuilt & custom models for fixed schemas	Extracts both explicit and inferred fields	Fully customizable via prompt engineering
Reasoning & Validation	No built-in reasoning	Built-in reasoning, validation, enrichment	Requires manual chaining and logic
Confidence & Grounding	Yes	Yes	No (requires custom implementation)
Ease of Use	Requires labeling/training for custom models	Zero-shot schema-based extraction, no labeling	Requires prompt tuning, orchestration, and engineering effort
Latency	Low	Medium	Variable (depends on deployment)
Scalability	Fully managed	Fully managed	Manual scaling required
Post-processing	Limited	Built-in rules for formatting, consistency	Fully user-defined
Multi-file Support	No	Yes (Pro Mode)	No (requires orchestration)
Knowledge Base Integration	No	Yes	Complex to implement
Security & Governance	Azure-compliant	Azure-compliant	Requires custom implementation
Total Cost of Ownership	Low for standard use cases	Moderate with built-in capabilities	High due to engineering and maintenance overhead

Demo Document Processing



Thank you

